

SlimX Reader: An Evidence-First Document Reader

Read, annotate, and ask grounded questions — entirely on your own machine.

1. What is SlimX Reader?

SlimX Reader is a local-first reading environment for PDF, Word, Markdown, text, and code documents. Unlike "chat with your PDF" tools, it is built around **evidence**: every answer it produces is grounded in retrieved passages from your documents, with visible page and section citations, and you can inspect exactly which context was used. Your files, annotations, and questions never leave your machine unless you explicitly enable a cloud provider.

2. Reading and annotation

The reader offers a professional PDF experience: zoom, page thumbnails, a document outline sidebar, and fast in-document search with match highlighting. Selecting any passage raises a small toolbar: pick one of five **highlight colors**, attach a **comment** that floats beside the text like a margin note, **ask** a question about the passage, or copy it. Clicking an existing highlight lets you recolor it, comment on it, or remove it. All annotations persist across reloads and can be exported to Markdown or JSON.

3. Grounded questions and answers

When a document is indexed, SlimX Reader sends it to **SlimX-RAG**, a local knowledge engine that parses, chunks, embeds, and indexes the content. Questions are answered by retrieving the most relevant chunks and passing them — clearly marked as untrusted context — to a local model such as Llama running under Ollama. The answer arrives with citations and a context-used panel, and any chunk or answer can be saved as an evidence note for later. If no model is available, retrieval and citations still work on their own.

4. Privacy by default

- Metadata lives in a local SQLite database; files stay on your filesystem.
- No telemetry, no analytics, no phone-home calls of any kind.
- Cloud providers are disabled by default and blocked at the API boundary.
- Secrets and tokens are redacted before anything is persisted or displayed.

5. Open source

SlimX Reader is MIT-licensed and developed in the open at github.com/slimx-ai/slimx-reader, alongside its sibling projects SlimX (model execution) and SlimX-RAG (retrieval). Contributions — issues, pull requests, and ideas — are welcome.

SlimX Reader — sample document for trying highlights, comments, indexing, and grounded Q&A.